

```
In [8]: 1 import math
        2 import numpy as np
        3 import pandas as pd
        4 import seaborn as sns
        5 import matplotlib.pyplot as plt
        6 import hashlib
```

```
In [9]: 1 # Importing Classifier Modules
        2 from sklearn.linear_model import LogisticRegression
        3 from sklearn.svm import SVC, LinearSVC
        4 from sklearn.neighbors import KNeighborsClassifier
        5 from sklearn.tree import DecisionTreeClassifier
        6 from sklearn.ensemble import RandomForestClassifier
        7 from sklearn.naive_bayes import GaussianNB
        8 from sklearn.linear_model import Perceptron
        9 from sklearn.linear_model import SGDClassifier
       10 from sklearn.neural_network import MLPClassifier
       11
       12 from sklearn.model_selection import learning_curve
       13 from sklearn.model_selection import train_test_split
       14 from sklearn.metrics import accuracy_score
       15 from sklearn.metrics import classification_report, confusion_matrix
       16 from sklearn import preprocessing
       17 from sklearn.model_selection import validation_curve
       18 from sklearn.model_selection import learning_curve
```

```

In [10]: 1 def plot_learning_curve(
2         estimator,
3         title,
4         X,
5         y,
6         axes=None,
7         ylim=None,
8         cv=None,
9         n_jobs=None,
10        train_sizes=np.linspace(0.1, 1.0, 5),
11    ):
12        """
13        Generate 3 plots: the test and training learning curve, the train
14        samples vs fit times curve, the fit times vs score curve.
15
16        Parameters
17        -----
18        estimator : estimator instance
19            An estimator instance implementing `fit` and `predict` methods
20            will be cloned for each validation.
21
22        title : str
23            Title for the chart.
24
25        X : array-like of shape (n_samples, n_features)
26            Training vector, where `n_samples` is the number of samples
27            and `n_features` is the number of features.
28
29        y : array-like of shape (n_samples) or (n_samples, n_features)
30            Target relative to `X` for classification or regression;
31            None for unsupervised learning.
32
33        axes : array-like of shape (3,), default=None
34            Axes to use for plotting the curves.
35
36        ylim : tuple of shape (2,), default=None
37            Defines minimum and maximum y-values plotted, e.g. (ymin, ymax)
38
39        cv : int, cross-validation generator or an iterable, default=None
40            Determines the cross-validation splitting strategy.
41            Possible inputs for cv are:
42
43            - None, to use the default 5-fold cross-validation,
44            - integer, to specify the number of folds.
45            - :term:`CV splitter`,
46            - An iterable yielding (train, test) splits as arrays of indices
47
48            For integer/None inputs, if `y` is binary or multiclass,
49            :class:`StratifiedKFold` used. If the estimator is not a classifier
50            or if `y` is neither binary nor multiclass, :class:`KFold` used.
51
52            Refer :ref:`User Guide <cross_validation>` for the various
53            cross-validators that can be used here.
54
55        n_jobs : int or None, default=None
56            Number of jobs to run in parallel.

```

```

57         ``None`` means 1 unless in a :obj:`joblib.parallel_backend`
58         ``-1`` means using all processors. See :term:`Glossary` <n_jobs>
59         for more details.
60
61     train_sizes : array-like of shape (n_ticks,)
62         Relative or absolute numbers of training examples that will
63         generate the learning curve. If the ``dtype`` is float, it is
64         as a fraction of the maximum size of the training set (that
65         determined by the selected validation method), i.e. it has to
66         be in the interval (0, 1]. Otherwise it is interpreted as absolute sizes of the
67         sets. Note that for classification the number of samples used
68         to be big enough to contain at least one sample from each class
69         (default: np.linspace(0.1, 1.0, 5))
70     """
71     if axes is None:
72         _, axes = plt.subplots(1, 3, figsize=(20, 5))
73
74     if ylim is not None:
75         axes[0].set_ylim(*ylim)
76     axes[0].set_xlabel("Training examples")
77     axes[0].set_ylabel("Score")
78
79     train_sizes, train_scores, test_scores, fit_times, _ = learning_
80         estimator,
81         X,
82         y,
83         cv=cv,
84         n_jobs=n_jobs,
85         train_sizes=train_sizes,
86         return_times=True,
87     )
88     train_scores_mean = np.mean(train_scores, axis=1)
89     train_scores_std = np.std(train_scores, axis=1)
90     test_scores_mean = np.mean(test_scores, axis=1)
91     test_scores_std = np.std(test_scores, axis=1)
92     fit_times_mean = np.mean(fit_times, axis=1)
93     fit_times_std = np.std(fit_times, axis=1)
94
95     # Plot learning curve
96     axes[0].grid()
97     axes[0].fill_between(
98         train_sizes,
99         train_scores_mean - train_scores_std,
100         train_scores_mean + train_scores_std,
101         alpha=0.1,
102         color="r",
103     )
104     axes[0].fill_between(
105         train_sizes,
106         test_scores_mean - test_scores_std,
107         test_scores_mean + test_scores_std,
108         alpha=0.1,
109         color="g",
110     )
111     axes[0].plot(
112         train_sizes, train_scores_mean, "o-", color="r", label="Traini
113

```

```
114 axes[0].plot(  
115     train_sizes, test_scores_mean, "o-", color="g", label="Cross  
116 )  
117 axes[0].legend(loc="best")  
118  
119 # Plot n_samples vs fit_times  
120 axes[1].grid()  
121 axes[1].plot(train_sizes, fit_times_mean, "o-")  
122 axes[1].fill_between(  
123     train_sizes,  
124     fit_times_mean - fit_times_std,  
125     fit_times_mean + fit_times_std,  
126     alpha=0.1,  
127 )  
128 axes[1].set_xlabel("Training examples")  
129 axes[1].set_ylabel("fit_times")  
130 axes[1].set_title("Scalability of the model")  
131  
132 # Plot fit_time vs score  
133 # Plot n_samples vs fit_times  
134 # Plot learning curve  
135 fit_time_argsort = fit_times_mean.argsort()  
136 fit_time_sorted = fit_times_mean[fit_time_argsort]  
137 test_scores_mean_sorted = test_scores_mean[fit_time_argsort]  
138 test_scores_std_sorted = test_scores_std[fit_time_argsort]  
139 axes[2].grid()  
140 axes[2].plot(fit_time_sorted, test_scores_mean_sorted, "o-")  
141 axes[2].fill_between(  
142     fit_time_sorted,  
143     test_scores_mean_sorted - test_scores_std_sorted,  
144     test_scores_mean_sorted + test_scores_std_sorted,  
145     alpha=0.1,  
146 )  
147 axes[2].set_xlabel("fit_times")  
148 axes[2].set_ylabel("Score")  
149 axes[2].set_title("Performance of the model")  
150  
151 return plt  
152
```

```
In [11]: 1 pokemon = pd.read_csv("pokemon.csv") # Dataset pokemon stats
          2 combats = pd.read_csv("combats.csv") # Datas:qxet pokemon battles and
          3
          4 pokemon["Type 2"] = pokemon["Type 2"].fillna("NA")
          5 pokemon.head()
          6
```

```
Out[11]:
```

	#	Name	Type 1	Type 2	HP	Attack	Defense	Sp. Atk	Sp. Def	Speed	Generation	Legendary
0	1	Bulbasaur	Grass	Poison	45	49	49	65	65	45	1	False
1	2	Ivysaur	Grass	Poison	60	62	63	80	80	60	1	False
2	3	Venusaur	Grass	Poison	80	82	83	100	100	80	1	False
3	4	Mega Venusaur	Grass	Poison	80	100	123	122	120	80	1	False
4	5	Charmander	Fire	NA	39	52	43	60	50	65	1	False

```
In [12]: 1 pokemon.columns
```

```
Out[12]: Index(['#', 'Name', 'Type 1', 'Type 2', 'HP', 'Attack', 'Defense', 'Sp.  
Atk',  
              'Sp. Def', 'Speed', 'Generation', 'Legendary'],  
              dtype='object')
```

## Data preprocessing

### First approach processing

We wanted to make some comparisons according to our beliefs of what is more impactfull in the outcome of a pokemon battle and for that we need a base comparison, we guided this first approach to an approach seen in [here \(https://github.com/kartikeya-rana/pokemon\\_battle/blob/master/Pokemon.ipynb\)](https://github.com/kartikeya-rana/pokemon_battle/blob/master/Pokemon.ipynb), in this approach the individual opted to create a new data set that consisted of pokemon 1 stats followed by pokemon 2 statistics and finnaly the winner of the battle, 1 in case the first pokemon wins 0 otherwise, for simplicity purposes it is not possible to have a match with no winner.

### Second approach processing

After that we thought about instead of concatenating both pokemons statistics we wanted to make a subtraction of the first pokemon with the second and compare the new results with the first ones, we also eliminated both the generation and the Legendary information, as well as the type information, since that we know that pokemons of different generation have the same overall raw power, and we also know that pokemon that are legendary do not have big boosts in their stats, and they do not win against every non legendary pokemon

## Third approach processing

This last approach had in account our knowledge of the domain, we believe that pokemons with higher speed and higher attack are more prone to win, therefore we extracted those values and included them in a new dataset.

```
In [13]: 1 data_one_hot_encoding = []
2 extra_data_one_hot_encoding = []
3 i = 0
4
5
6 # for each tuple of combats.csv
7 for t in combats.itertuples():
8     i += 1
9     first_pokemon = t[1] # get the first pokemon
10    second_pokemon = t[2] # get the second pokemon
11    winner = t[3]         # get the winner
12
13    if i <= 5001:
14        x = pokemon.loc[pokemon["#"]==first_pokemon].values[:, 2:][0]
15        y = pokemon.loc[pokemon["#"]==second_pokemon].values[:, 2:][0]
16        #diff = (x-y)[:6] # difference between "base stats hp...."
17        z = np.concatenate((x,y))
18        if winner == first_pokemon:
19            z = np.append(z, [0])
20        else:
21            z = np.append(z, [1])
22
23
24        data_one_hot_encoding.append(z)
25    elif i < 10000:
26        x = pokemon.loc[pokemon["#"]==first_pokemon].values[:, 2:][0]
27        y = pokemon.loc[pokemon["#"]==second_pokemon].values[:, 2:][0]
28        #diff = (x-y)[:6] # difference between "base stats hp...."
29        z = np.concatenate((x,y))
30        if winner == first_pokemon:
31            z = np.append(z, [0])
32        else:
33            z = np.append(z, [1])
34
35
36        extra_data_one_hot_encoding.append(z)
37    else:
38        break
39
40
41
42 data_one_hot_encoding = np.asarray(data_one_hot_encoding)
43
```

```

In [14]: 1 data_diff_base_stats = []
2 extra_data_diff_base_stats = []
3 i = 0
4
5
6 # for each tuple of combats.csv
7 for t in combats.itertuples():
8     i += 1
9     first_pokemon = t[1] # get the first pokemon
10    second_pokemon = t[2] # get the second pokemon
11    winner = t[3]         # get the winner
12
13    if i <= 5001:
14        x = pokemon.loc[pokemon["#"]==first_pokemon].values[:, 2:][0]
15        y = pokemon.loc[pokemon["#"]==second_pokemon].values[:, 2:][0]
16        diff = (x[2:8]-y[2:8]) # difference between "base stats hp...
17
18        z = []
19        z = np.append(diff,z,0)
20
21
22        if winner == first_pokemon:
23            z = np.append(z, [0])
24        else:
25            z = np.append(z, [1])
26
27        data_diff_base_stats.append(z)
28
29    elif i < 10000:
30        x = pokemon.loc[pokemon["#"]==first_pokemon].values[:, 2:][0]
31        y = pokemon.loc[pokemon["#"]==second_pokemon].values[:, 2:][0]
32        diff = (x[2:8]-y[2:8]) # difference between "base stats hp...
33
34        z = []
35        z = np.append(diff,z,0)
36
37
38        if winner == first_pokemon:
39            z = np.append(z, [0])
40        else:
41            z = np.append(z, [1])
42
43        extra_data_diff_base_stats.append(z)
44    else:
45        break
46
47    data_diff_base_stats = np.asarray(data_diff_base_stats)
48    data_diff_base_stats = data_diff_base_stats[:, :-1].astype(int)
49
50    extra_data_diff_base_stats = np.asarray(extra_data_diff_base_stats)
51
52

```

```
In [15]: 1 pokemon_with_types = pokemon.copy()
2 types_list = [x for x in pokemon["Type 2"].unique()] + [x for x in po
3
4 types_list = list(set(types_list))
5
6 print(types_list.remove("NA"))
7 # make NA -> 0
8 types_list = ["NA"] + types_list
9
10 types_map = { x : types_list.index(x) for x in types_list}
11
12 pokemon_with_types["Type 1"] = pokemon_with_types["Type 1"].map(types
13 pokemon_with_types["Type 2"] = pokemon_with_types["Type 2"].map(types
14
15 pokemon.head()
16
17 print(types_map)
18 print(types_list)
```

None

```
{'NA': 0, 'Water': 1, 'Psychic': 2, 'Poison': 3, 'Bug': 4, 'Ghost': 5,
'Electric': 6, 'Fire': 7, 'Steel': 8, 'Flying': 9, 'Rock': 10, 'Normal':
11, 'Fighting': 12, 'Ground': 13, 'Dragon': 14, 'Fairy': 15, 'Dark': 16,
'Grass': 17, 'Ice': 18}
['NA', 'Water', 'Psychic', 'Poison', 'Bug', 'Ghost', 'Electric', 'Fire',
'Steel', 'Flying', 'Rock', 'Normal', 'Fighting', 'Ground', 'Dragon', 'Fa
iry', 'Dark', 'Grass', 'Ice']
```



```

In [16]: 1 data_with_types = []
2 extra_data_with_types = []
3 i = 0
4
5
6 # for each tuple of combats.csv
7 for t in combats.itertuples():
8     i += 1
9     first_pokemon = t[1] # get the first pokemon
10    second_pokemon = t[2] # get the second pokemon
11    winner = t[3]         # get the winner
12
13    if i <= 5001:
14        x = pokemon_with_types.loc[pokemon["#"]==first_pokemon]
15        x = x.drop(columns=["Name", "#"]).values[0]
16        y = pokemon_with_types.loc[pokemon["#"]==second_pokemon].drop
17
18
19        z = np.concatenate((x,y))
20
21        if winner == first_pokemon:
22            z = np.append(z, [0])
23        else:
24            z = np.append(z, [1])
25
26        data_with_types.append(z)
27    elif i < 10000:
28        x = pokemon_with_types.loc[pokemon["#"]==first_pokemon]
29        x = x.drop(columns=["Name", "#"]).values[0]
30        y = pokemon_with_types.loc[pokemon["#"]==second_pokemon].drop
31
32
33        z = np.concatenate((x,y))
34
35        if winner == first_pokemon:
36            z = np.append(z, [0])
37        else:
38            z = np.append(z, [1])
39
40        extra_data_with_types.append(z)
41    else:
42        break
43
44
45 data_with_types = np.asarray(data_with_types)
46 data_with_types= pd.DataFrame(data_with_types, columns=[
47     'Type 1', 'Type 2', 'HP', 'Attack', 'Defense', 'Sp. Atk',
48     'Sp. Def', 'Speed', 'Generation', 'Legendary',
49     'Type 1', 'Type 2', 'HP', 'Attack', 'Defense', 'Sp. Atk',
50     'Sp. Def', 'Speed', 'Generation', 'Legendary', 'Win'] )
51
52 data_with_types = data_with_types.astype({
53     'Type 1':'int32', 'Type 2':'int32', 'HP':'int32', 'Attack':'in
54     'Sp. Def':'int32', 'Speed':'int32', 'Generation':'int32', 'Leg
55     'Type 1':'int32', 'Type 2':'int32', 'HP':'int32', 'Attack':'in
56     'Sp. Def':'int32', 'Speed':'int32', 'Generation':'int32', 'Leg

```

```
57
58 extra_data_with_types = np.asarray(extra_data_with_types)
59 extra_data_with_types= pd.DataFrame(extra_data_with_types, columns=[
60     'Type 1', 'Type 2', 'HP', 'Attack', 'Defense', 'Sp. Atk',
61     'Sp. Def', 'Speed', 'Generation', 'Legendary' ,
62     'Type 1', 'Type 2', 'HP', 'Attack', 'Defense', 'Sp. Atk',
63     'Sp. Def', 'Speed', 'Generation', 'Legendary', 'Win'] )
64
65 extra_data_with_types = extra_data_with_types.astype({
66     'Type 1': 'int32', 'Type 2': 'int32', 'HP': 'int32', 'Attack': 'in
67     'Sp. Def': 'int32', 'Speed': 'int32', 'Generation': 'int32', 'Leg
68     'Type 1': 'int32', 'Type 2': 'int32', 'HP': 'int32', 'Attack': 'in
69     'Sp. Def': 'int32', 'Speed': 'int32', 'Generation': 'int32', 'Leg
70
71
72
73 X_imp_feat = data_with_types[['Type 1', 'Type 2', 'Attack', 'Speed']]
74
75 y = data_with_types["Win"]
76
77 print(y)
78
79
80
```

```
0      1
1      1
2      1
3      1
4      0
..
4996   0
4997   1
4998   1
4999   0
5000   1
Name: Win, Length: 5001, dtype: int32
```

```
In [17]: 1 extra_data_with_types.head()
```

Out[17]:

	Type 1	Type 2	HP	Attack	Defense	Sp. Atk	Sp. Def	Speed	Generation	Legendary	...	Type 2	HP	At
0	8	0	40	55	70	45	60	30	5	0	...	0	38	
1	5	17	85	110	76	65	82	56	6	0	...	0	30	
2	18	9	90	85	100	95	125	85	1	1	...	1	110	
3	11	0	70	70	70	70	70	70	3	0	...	0	95	
4	11	9	85	120	70	50	60	100	4	0	...	3	40	

5 rows × 21 columns

Data correlation number of wins

number of wins is correlated with speed and attack being special attack and defense are also important factors as well as being legendary Contrary to my beliefs generations are quite well balanced, since there is almost no correlation between the 2 We can also see that the base stats do not have much correlation with generation, therefore pokemons that belong to different generations are on an equal foot in terms of base power.


## Data correlation legendary or not

Once again through this correlation heatmap it is possible to say that generation has no important role to classify a pokemon as legendary, one thing that strikes out is Sp. Attack and Attack as well as Sp. Defense this seems to be somehow correlated to classifying a pokemon as legendary or not

```
In [18]: 1 unique_ids=pokemon['#']
2 wins_by_id = []
3 for _id in unique_ids:
4     wins_by_id.append([_id ,(combats["Winner"] == _id).sum() ])
5
6 wins_by_id = np.asarray(wins_by_id)
7
8 pokemon["Wins"] = wins_by_id[:,1]
9 pokemon["Legendary"] = pokemon["Legendary"].astype(int)
10
11
12
```

```
In [19]: 1
          2 plt.figure(figsize=(20,8))
          3 sns.heatmap(pokemon.corr(),square=True, vmin=-1, vmax=1, annot=True)
```

Out[19]: <AxesSubplot:>



## Data correlation number of wins

number of wins is correlated with speed and attack being special attack and defense are also important factors as well as being legendary Contrary to my beliefs generations are quite well balanced, since there is almos no correlation between the 2 We can also see that the base stats do not have much correlation with generation, therefore pokemons that belong to diferent generations are on an equal foot in terms of base power.

## Data correlation legendary or not

Once again through this correlation heatmap it is possible to say that generation as no important role to classify a pokemon as legendary, one thing that strikes out is Sp. Attack and Attack as well as Sp. Defense this seem to be somehow correlated to classifying a pokemon as legendary or not

## Graph that shows the correlation between number of wins and speed


Overall the speed of the pokemon is more impactfull in the outcome of battle than most of the other attributes.

```


In [20]: 1 #boxplot of Attack vs. Legendary
2 plt.figure(figsize=(8, 4))
3 sns.boxplot(x='Legendary',y='Wins',data=pokemon, palette='rainbow')
4
5 #stripplot of speed Wins relation
6 plt.figure(figsize=(20,8))
7 speed_wins = pokemon[['Speed', 'Wins']].groupby(['Speed'], as_index=False)
8 speed_wins.sort_values(by='Speed',ascending=True).plot(kind='line')
9 #stripplot of Attack vs. Legendary, palette='rainbow'
10 plt.figure(figsize=(20,8))
11 speed_wins = pokemon[['Attack', 'Wins']].groupby(['Attack'], as_index=False)
12 speed_wins.sort_values(by='Attack',ascending=True).plot(kind='line')
13

```


Out[20]: <AxesSubplot:xlabel='Attack'>



<Figure size 1440x576 with 0 Axes>



<Figure size 1440x576 with 0 Axes>





**1 type pokemon go brr cause strong**

```

In [21]: 1 type_1 = pokemon[['Type 1', 'Wins']].groupby(['Type 1'], as_index=False)
          2 type_1.sort_values(by='Wins', ascending=False).plot(kind='bar')
          3
          4 type_2 = pokemon[['Type 2', 'Wins']].groupby(['Type 2'], as_index=False)
          5 type_2.sort_values(by='Wins', ascending=False).plot(kind='bar')

```


Out[21]: <AxesSubplot:xlabel='Type 2'>






```
In [22]: 1 plt.figure(figsize=(48, 12))
2 plt.xlabel("Name")
3 plt.ylabel("Wins")
4 plt.bar(pokemon['Name'][:20], pokemon['Wins'][:20])
```

Out[22]: <BarContainer object of 20 artists>




```
In [23]: 1
2 type_1 = pokemon[['Attack', 'Wins']].groupby(['Attack'], as_index=False)
3 type_1.sort_values(by='Wins', ascending=True).plot(kind='bar', figsize
4 #Number of wins is directly proportional too attack damage
```

Out[23]: Text(0, 0.5, 'Wins')



```
In [24]: 1 type_1 = pokemon[['Speed', 'Wins']].groupby(['Speed'], as_index=False)
2 type_1.sort_values(by='Wins', ascending=True).plot(kind='bar', figsize=
3 #Number of wins is directly proportional too Speed
4 type_1 = pokemon[['Speed', 'Wins']].groupby(['Speed'], as_index=False)
5 type_1.sort_values(by='Speed', ascending=True).plot(kind='line', figsi
6 #Number of wins is directly proportional too Speed
```

Out[24]: <AxesSubplot: xlabel='Speed'>



## Classifier Results Not Having in Account Types

As we know types take an important role in pokemon battles, for example a electric pokemon deals 2 times more damage to an water pokemon than to a fire pokemon and to times less damage to a rock pokemon, this interactions might prove challenging to some classification algorithms.

## Second approach

Features usados HP Diff int64 Attack Diff int64 Defens Diff int64 Sp. Atk diff int64 Sp. Def dif int64 Speed diff int64

```
In [25]: 1 data_pd_frame = pd.DataFrame(data_diff_base_stats, columns=["HP Diff"
2 # "HP Diff & Attack Diff & Defens Diff & Sp. Atk diff & Sp. Def dif &
3 # & -20 & -6 & 10 & -15 & 10 & -19
4 for column in data_pd_frame.columns:
5     data_pd_frame[column] = data_pd_frame[column].astype(int)
6
7
8 extra_data_pd_frame = pd.DataFrame(extra_data_diff_base_stats, column
9 # "HP Diff & Attack Diff & Defens Diff & Sp. Atk diff & Sp. Def dif &
10 # & -20 & -6 & 10 & -15 & 10 & -19
11 for column in data_pd_frame.columns:
12     extra_data_pd_frame[column] = extra_data_pd_frame[column].astype(int)
```

```
In [26]: 1 data_pd_frame.dtypes
2 data_pd_frame.head(1)
```

```
Out[26]:
```


	HP Diff	Attack Diff	Defens Diff	Sp. Atk diff	Sp. Def dif	Speed diff
0	-20	-6	10	-15	10	-19

```
In [27]: 1 X_diff_stats = data_pd_frame.values
2
3 X_diff_train, X_diff_test, y_train, y_test = train_test_split(X_diff_stats,
```

```
In [28]: 1 # Logistic Regression
2 clf = LogisticRegression(max_iter=10000)
3 clf.fit(X_diff_train, y_train)
4 y_pred_log_reg = clf.predict(X_diff_test)
5 acc_log_reg = round( clf.score(X_diff_test, y_test) * 100, 2)
6 print(str(acc_log_reg) + ' percent')
7 sns.heatmap(confusion_matrix(y_test, y_pred_log_reg), square=True, an
```


88.41 percent

```
Out[28]: <AxesSubplot:>
```



```
In [29]: 1 plot_learning_curve(
2         clf, "Logistic Regression", X_diff_stats, y, ylim=(-1, 1.01), cv=5
3     )
```

Out[29]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



```
In [30]: 1 pd.DataFrame(zip(data_pd_frame.columns, np.transpose(clf.coef_[0])),
```

Out[30]:


	features	coef
0	HP Diff	-0.000941
1	Attack Diff	-0.009897
2	Defens Diff	-0.003292
3	Sp. Atk diff	0.002448
4	Sp. Def dif	-0.001284
5	Speed diff	-0.061518

```
In [31]: 1 # Random Forest Classifier difference between most imp features
2 # Every diff
3 clf = RandomForestClassifier(n_estimators=100)
4 model = clf.fit(X_diff_train,y_train)
5 pred = model.predict(X_diff_test)
6 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
7 print('Accuracy :', accuracy_score(pred, y_test))
8 print(classification_report(y_test, pred))
9 sns.heatmap(confusion_matrix(y_test, pred), square=True, annot=True)
```

```
Accuracy : 0.9520383693045563
```


	precision	recall	f1-score	support
0	0.94	0.95	0.95	574
1	0.96	0.95	0.96	677
accuracy			0.95	1251
macro avg	0.95	0.95	0.95	1251
weighted avg	0.95	0.95	0.95	1251

Out[31]: <AxesSubplot:>



```
In [32]: 1 plot_learning_curve(
2         clf, "Random Forest Classifier", X_diff_stats, y, ylim=(-1, 1.01),
3     )
```

Out[32]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>




# Leave Only Attack and Speed diff

```
In [33]: 1 X_Most_Imp_Features = data_pd_frame.drop(columns=["HP Diff", "Sp. Atk
2 X_train, X_test, y_train, y_test = train_test_split(X_Most_Imp_Featur
3 randomFlorestclf = RandomForestClassifier(n_estimators=100)
4 model = randomFlorestclf.fit(X_train,y_train)
5 pred = model.predict(X_test)
6 # print('Accuracy of {}'.format(name), accuracy_score(pred, y_test))
7 print('Accuracy :', accuracy_score(pred, y_test))
8 print(classification_report(y_test, pred))
9 sns.heatmap(confusion_matrix(y_test,pred), square=True, annot=True)
```

```
Accuracy : 0.9368505195843325
```

	precision	recall	f1-score	support
0	0.93	0.93	0.93	574
1	0.94	0.94	0.94	677
accuracy			0.94	1251
macro avg	0.94	0.94	0.94	1251
weighted avg	0.94	0.94	0.94	1251

Out[33]: <AxesSubplot:>




```
In [34]: 1 X_Most_Imp_Features.head(2)
```

Out[34]:

	Attack Diff	Speed diff
0	-6	-19
1	-39	0

```
In [35]: 1 plot_learning_curve(
          2     randomFlorestclf, "Random Forest Classifier", X_Most_Imp_Features
          3 )
```

Out[35]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



```
In [36]: 1 X_Most_Imp_Features.head(3)
```


Out[36]:

	Attack Diff	Speed diff
0	-6	-19
1	-39	0
2	-35	0

```
In [37]: 1 # Logistic Regression difference between most imp features
2 # Attack Diff Speed Dif
3 clf = LogisticRegression(C=100,max_iter=10000)
4 clf.fit(X_train, y_train)
5 y_pred_log_reg = clf.predict(X_test)
6 acc_log_reg = round( clf.score(X_test, y_test) * 100, 2)
7 print(str(acc_log_reg) + ' percent')
8 sns.heatmap(confusion_matrix(y_test,y_pred_log_reg), square=True, ann
```


88.73 percent

Out[37]: <AxesSubplot:>



```
In [38]: 1 plot_learning_curve(
2         clf, "Logistic Regression", X_Most_Imp_Features, y,ylim=(-1, 1.01
3         )
```

Out[38]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



## First approach

Pokemon stats in the same row types features encoded as a mapping

```
In [39]: 1 data = data_with_types.drop(columns=["Win", "Legendary", "Generation"])
```



In [40]: 1 data.head(3)

Out[40]:

	Type_1	Type_2	HP	Attack	Defense	Sp. Atk	Sp. Def	Speed	Type_1	Type_2	HP	Attack	Defense	Sp. Atk	Sp. Def	Speed
0	10	13	50	64	50	45	50	41	17	16	70	70	40	60	50	41
1	17	12	91	90	72	90	129	108	10	12	91	129	90	72	90	108
2	15	9	55	40	85	80	105	40	2	0	75	75	75	125	80	105


In [41]:

```

1 X = data.values[:, :-1].astype(int)
2
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
4
5 clf = LogisticRegression(C=0.1,max_iter=10000)
6 clf.fit(X_train, y_train)
7 y_pred_log_reg = clf.predict(X_test)
8 acc_log_reg = round( clf.score(X_test, y_test) * 100, 2)
9 print(str(acc_log_reg) + ' percent')
10 sns.heatmap(confusion_matrix(y_test,y_pred_log_reg), square=True, annot=True, cbar=True)
```


76.66 percent

Out[41]: <AxesSubplot:>



```
In [42]: 1 plot_learning_curve(
2         clf, "Logistic Regression", X, y, ylim=(-1, 1.01), cv=5, n_jobs=4
3     )
```

Out[42]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



```
In [43]: 1 print(X[0])
```


[10 13 50 64 50 45 50 41 17 16 70 70 40 60 40]

```
In [44]: 1 # Random forest Classifier with every stat of both pokemons including
2 clf = RandomForestClassifier(n_estimators=100)
3 model = clf.fit(X_train,y_train)
4 pred = model.predict(X_test)
5 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
6 print('Accuracy :', accuracy_score(pred, y_test))
7 print(classification_report(y_test, pred))
8 sns.heatmap(confusion_matrix(y_test, pred), square=True, annot=True)
```

```
Accuracy : 0.7929656274980016
```


	precision	recall	f1-score	support
0	0.78	0.77	0.77	574
1	0.81	0.81	0.81	677
accuracy			0.79	1251
macro avg	0.79	0.79	0.79	1251
weighted avg	0.79	0.79	0.79	1251

Out[44]: <AxesSubplot:>



```
In [45]: 1 plot_learning_curve(
2         clf, "Random Forest Classifier", X, y, ylim=(-1, 1.01), cv=5, n_jobs=
3         )
```

Out[45]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



```


In [46]: 1 # MLP classifier with every stat as well as types map
2 clf = MLPClassifier(solver='adam', alpha=1e-5, hidden_layer_sizes=(10
3 model = clf.fit(X_train, y_train)
4 pred = model.predict(X_test)
5 print('Accuracy :', accuracy_score(pred, y_test))
6 print(classification_report(y_test, pred))
7 sns.heatmap(confusion_matrix(y_test, pred), square=True, annot=True)
8

```

Accuracy : 0.7513988808952837


	precision	recall	f1-score	support
0	0.80	0.61	0.69	574
1	0.72	0.87	0.79	677
accuracy			0.75	1251
macro avg	0.76	0.74	0.74	1251
weighted avg	0.76	0.75	0.75	1251

Out[46]: <AxesSubplot:>



```
In [47]: 1 plot_learning_curve(  
2         clf, "MLP Classifier", X, y, ylim=(-1, 1.01), cv=5, n_jobs=4  
3     )
```

Out[47]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



**All base stats normally appended to the data types as one hot encoding**

```

In [48]: 1 # limit to categorical data using df.select_dtypes()
2 data_one_hot_encoding = pd.DataFrame( data_one_hot_encoding, columns=[
3                                     'Defense', 'Sp. At
4                                     'Generation', 'Leg
5                                     'HP', 'Attack',
6                                     'Defense', 'Sp. At
7                                     'Generation', 'Leg
8 save = data_one_hot_encoding.copy()
9
10 data_one_hot_encoding = data_one_hot_encoding.astype({'HP' : 'int32',
11                                                         'Defense': 'int
12                                                         'Sp. Def' : 'in
13                                                         'Generation': '
14 print(data_one_hot_encoding.dtypes)
15 X_objects = data_one_hot_encoding.select_dtypes(include=[object])
16 X_objects.head(3)

```

```

Type 1      object
Type 2      object
HP          int32
Attack      int32
Defense     int32
Sp. Atk     int32
Sp. Def     int32
Speed       int32
Generation  int32
Legendary   int32
Type 1      object
Type 2      object
HP          int32
Attack      int32
Defense     int32
Sp. Atk     int32
Sp. Def     int32
Speed       int32
Generation  int32
Legendary   int32
Winner      int32
dtype: object

```

```

Out[48]:

```

	Type 1	Type 2	Type 1	Type 2
0	Rock	Ground	Grass	Dark
1	Grass	Fighting	Rock	Fighting
2	Fairy	Flying	Psychic	NA

In [49]:

```

1  # 1. INSTANTIATE
2  # encode labels with value between 0 and n_classes-1.
3  le = preprocessing.LabelEncoder()
4
5
6  # 2/3. FIT AND TRANSFORM
7  # use df.apply() to apply le.fit_transform to all columns
8  X_objects_fitted = X_objects.apply(le.fit_transform)
9  X_objects_fitted.head()
10
11 # limit to categorical data using df.select_dtypes()
12 enc = preprocessing.OneHotEncoder()
13
14 # 2. FIT
15 enc.fit(X_objects_fitted)
16
17 # 3. Transform
18 onehotlabels = enc.transform(X_objects_fitted).toarray()
19 onehotlabels.shape
20 panda_Hot = pd.DataFrame(onehotlabels)
21 X_data_one_Hot = data_one_hot_encoding.drop(columns=["Winner", "Type"]
22
23 X_data_one_Hot = X_data_one_Hot.join(panda_Hot)
24 X_data_one_Hot.head()

```

Out[49]:

	HP	Attack	Defense	Sp. Atk	Sp. Def	Speed	Generation	Legendary	HP	Attack	...	64	65	66
0	50	64	50	45	50	41	2	0	70	70	...	0.0	0.0	0.0
1	91	90	72	90	129	108	5	1	91	129	...	0.0	0.0	0.0
2	55	40	85	80	105	40	2	0	75	75	...	0.0	0.0	0.0
3	40	40	40	70	40	20	2	0	77	120	...	0.0	0.0	0.0
4	70	60	125	115	70	55	1	0	20	10	...	0.0	0.0	0.0

5 rows × 90 columns

In [50]:

```

1 X_train, X_test, y_train, y_test = train_test_split(X_data_one_Hot, y
2 print(X_train.shape)
3 print(y_train.shape)

```

(3750, 90)

(3750, )

```

In [51]: 1 #this is bad3!
2 clf = MLPClassifier(solver='adam', alpha=1e-5, hidden_layer_sizes=(10
3
4 y = data_one_hot_encoding.Winner
5 print(y.shape)
6
7 model = clf.fit(X_train, y_train)
8 pred = model.predict(X_test)
9 print(set(pred))
10 print(set(y_test))
11 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
12 print('Accuracy :', accuracy_score(pred, y_test))
13 print(classification_report(y_test, pred))

```

(5001,)

{0, 1}

{0, 1}

Accuracy : 0.9072741806554756


	precision	recall	f1-score	support
0	0.92	0.88	0.90	574
1	0.90	0.93	0.92	677
accuracy			0.91	1251
macro avg	0.91	0.91	0.91	1251
weighted avg	0.91	0.91	0.91	1251

```

In [52]: 1 plot_learning_curve(
2         clf, "MLP Classifier", X_data_one_Hot, y, ylim=(-1, 1.01), cv=5, n
3         )

```

Out[52]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>





```

In [53]: 1 # Random Forest Classifier difference between most imp features
          2 # Attack Diff Speed Dif
          3 clf = RandomForestClassifier(n_estimators=100)
          4 model = clf.fit(X_train,y_train)
          5
          6 pred = model.predict(X_test)
          7 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
          8 print('Accuracy :', accuracy_score(pred, y_test))
          9 print(classification_report(y_test, pred))
         10 sns.heatmap(confusion_matrix(y_test, pred), square=True, annot=True)

```

```


Accuracy : 0.8808952837729817
              precision    recall  f1-score   support

             0       0.87       0.87       0.87       574
             1       0.89       0.89       0.89       677

 accuracy
macro avg       0.88       0.88       0.88       1251
weighted avg    0.88       0.88       0.88       1251


```

Out[53]: <AxesSubplot:>



```
In [54]: 1 plot_learning_curve(
2         clf, "Random Forest Classifier", X_data_one_Hot, y, ylim=(-1, 1.01)
3     )
```


Out[54]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



```
In [55]: 1 clf = LogisticRegression(C=0.1,max_iter=10000)
2         clf.fit(X_train, y_train)
3         y_pred_log_reg = clf.predict(X_test)
4         acc_log_reg = round( clf.score(X_test, y_test) * 100, 2)
5         print(str(acc_log_reg) + ' percent')
6         sns.heatmap(confusion_matrix(y_test,y_pred_log_reg ), square=True, an
```


88.33 percent

Out[55]: <AxesSubplot:>



```
In [56]: 1 plot_learning_curve(
2         clf, "Logistic Regression", X_data_one_Hot, y, ylim=(-1, 1.01), cv
3         )
```

Out[56]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



## Third Approach

Use only attack speed and types as a map

```
In [57]: 1 X = data_with_types[["Type 1", "Type 2", "Attack", "Speed"]]
2         X.head(3)
```

Out[57]:

	Type 1	Type 1	Type 2	Type 2	Attack	Attack	Speed	Speed
0	10	17	13	16	64	70	41	60
1	17	10	12	12	90	129	108	108
2	15	2	9	0	40	75	40	40

```
In [58]: 1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
2         print(X_train.shape)
3         print(y_train.shape)
```

(3750, 8)  
(3750,)


```
In [59]: 1 #this is bad3!
2 clf = MLPClassifier(solver='adam', alpha=1e-5, hidden_layer_sizes=(10
3
4 y = data_one_hot_encoding.Winner
5 print(y.shape)
6
7 model = clf.fit(X_train, y_train)
8 pred = model.predict(X_test)
9 print(set(pred))
10 print(set(y_test))
11 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
12 print('Accuracy :', accuracy_score(pred, y_test))
13 print(classification_report(y_test, pred))
```

```
(5001,)
{0, 1}
{0, 1}
Accuracy : 0.9232613908872902
```

	precision	recall	f1-score	support
0	0.95	0.88	0.91	574
1	0.90	0.96	0.93	677
accuracy			0.92	1251
macro avg	0.93	0.92	0.92	1251
weighted avg	0.92	0.92	0.92	1251

```
In [60]: 1 plot_learning_curve(
2         clf, "MLP Classifier", X, y, ylim=(-1, 1.01), cv=5, n_jobs=4
3     )
```

```
Out[60]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>
```



```


In [61]: 1 # Random Forest Classifier difference between most imp features
          2 # Attack Diff Speed Dif
          3 clf = RandomForestClassifier(n_estimators=100)
          4 model = clf.fit(X_train,y_train)
          5
          6 pred = model.predict(X_test)
          7 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
          8 print('Accuracy :', accuracy_score(pred, y_test))
          9 print(classification_report(y_test, pred))
         10 sns.heatmap(confusion_matrix(y_test, pred), square=True, annot=True)

```

Accuracy : 0.9272581934452439


	precision	recall	f1-score	support
0	0.92	0.93	0.92	574
1	0.94	0.93	0.93	677
accuracy			0.93	1251
macro avg	0.93	0.93	0.93	1251
weighted avg	0.93	0.93	0.93	1251

Out[61]: <AxesSubplot:>



```
In [62]: 1 plot_learning_curve(
2         clf, "Random Forest Classifier", X, y, ylim=(-1, 1.01), cv=5, n_jobs=
3         )
```


Out[62]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



```
In [63]: 1
2 clf = LogisticRegression(C=0.1,max_iter=10000)
3 clf.fit(X_train, y_train)
4 y_pred_log_reg = clf.predict(X_test)
5 acc_log_reg = round( clf.score(X_test, y_test) * 100, 2)
6 print(str(acc_log_reg) + ' percent')
7 sns.heatmap(confusion_matrix(y_test,y_pred_log_reg ), square=True, a
```


88.49 percent

Out[63]: <AxesSubplot:>



```
In [64]: 1 plot_learning_curve(
2         clf, "Logistic Regression", X_data_one_Hot, y, ylim=(-1, 1.01), cv
3         )
```

Out[64]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



## Use difference stats and types mapped to integers

```
In [65]: 1 X_Most_Imp_Features = X_Most_Imp_Features.join(data_with_types[["Type
2         ]
```

```
In [66]: 1 X_Most_Imp_Features.head(3)
```

Out[66]:

	Attack Diff	Speed diff	Type 1	Type 1	Type 2	Type 2
0	-6	-19	10	17	13	16
1	-39	0	17	10	12	12
2	-35	0	15	2	9	0


```
In [83]: 1 X_train, X_test, y_train, y_test = train_test_split(X_Most_Imp_Feature
2 randomFlorestclf = RandomForestClassifier(n_estimators=100)
3 model = randomFlorestclf.fit(X_train,y_train)
4 pred = model.predict(X_test)
5 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
6 print('Accuracy :', accuracy_score(pred, y_test))
7 print(classification_report(y_test, pred))
```

```
Accuracy : 0.9496402877697842
```

	precision	recall	f1-score	support
0	0.94	0.95	0.95	574
1	0.96	0.95	0.95	677
accuracy			0.95	1251
macro avg	0.95	0.95	0.95	1251
weighted avg	0.95	0.95	0.95	1251


```
In [84]: 1 plot_learning_curve(
2         randomFlorestclf, "Random Forest Classifier", X_Most_Imp_Features
3     )
```

Out[84]: <module 'matplotlib.pyplot' from '/opt/anaconda/lib/python3.8/site-packages/matplotlib/pyplot.py'>



```
In [85]: 1 sns.heatmap(confusion_matrix(y_test, pred), square=True, annot=True)
```

Out[85]: <AxesSubplot:>





```
In [70]: 1
          2 train_features, test_features, train_labels, test_labels = train_test
          3
```

```
In [ ]: 1 from sklearn.model_selection import RandomizedSearchCV
          2
          3 # Number of trees in random forest
          4 n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000,
          5 # Number of features to consider at every split
          6 max_features = ['auto', 'sqrt'])
          7 # Maximum number of levels in tree
          8 max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
          9 max_depth.append(None)
         10 # Minimum number of samples required to split a node
         11 min_samples_split = [2, 5, 10]
         12 # Minimum number of samples required at each leaf node
         13 min_samples_leaf = [1, 2, 4]
         14 # Method of selecting samples for training each tree
         15 bootstrap = [True, False]
         16
         17 # Create the random grid
         18 random_grid = {'n_estimators': n_estimators,
         19                 'max_features': max_features,
         20                 'max_depth': max_depth,
         21                 'min_samples_split': min_samples_split,
         22                 'min_samples_leaf': min_samples_leaf,
         23                 'bootstrap': bootstrap}
         24
         25 # Use the random grid to search for best hyperparameters
         26 # First create the base model to tune
         27 rf = RandomForestClassifier(random_state = 42)
         28 # Random search of parameters, using 3 fold cross validation,
         29 # search across 100 different combinations, and use all available cor
         30 rf_random = RandomizedSearchCV(estimator=rf, param_distributions=rand
         31                               n_iter = 10000, scoring='neg_mean_abso
         32                               cv = 3, verbose=2, random_state=42, n_j
         33                               return_train_score=True)
         34
         35
         36
         37 # Fit the random search model
         38 rf_random.fit(train_features , train_labels);
```

## Evaluate the Default Model

```
In [73]: 1 base_model = RandomForestClassifier(n_estimators = 100, random_state
2         base_model.fit(train_features, train_labels)
3         pred = base_model.predict(test_features)
4
5         base_acc = accuracy_score(pred, test_labels)
6
7         # print('Accuracy of {}'.format(name), accuracy_score(pred, y_test))
8         print('Accuracy :', accuracy_score(pred, test_labels))
9         print(classification_report(pred, test_labels))
10
```

Accuracy : 0.947242206235012

	precision	recall	f1-score	support
0	0.95	0.94	0.94	580
1	0.95	0.96	0.95	671
accuracy			0.95	1251
macro avg	0.95	0.95	0.95	1251
weighted avg	0.95	0.95	0.95	1251

## Evaluate the Best Random Search Model

```
In [ ]: 1 best_random = rf_random.best_estimator_
2
3         pred = best_random.predict(test_features)
4
5         # print('Accuracy of {}'.format(name), accuracy_score(pred, y_test))
6         random_acc = accuracy_score(pred, test_labels)
7         print('Accuracy :', accuracy_score(pred, test_labels))
8         print(classification_report(pred, test_labels))
9
10        print('Improvement of {:.2f}%'.format( 100 * (random_acc - base_acc))
```

## Introducing new data to see how the final model behaves

```

In [81]: 1
2 test_labels = extra_data_pd_frame ["Win"].astype(int).values
3
4 X_new = extra_data_pd_frame[["Attack Diff", "Speed diff"]]
5 X_new = X_new.join(extra_data_with_types[["Type 1", "Type 2"]])
6
7 pred = randomFlorestclf.predict(X_new)
8
9 # print('Accuracy of {}:'.format(name), accuracy_score(pred, y_test))
10 random_acc = accuracy_score(pred, test_labels)
11 print('Accuracy :', accuracy_score(pred, test_labels))
12 print(classification_report(pred, test_labels))
13
14
15
16
17
18
19 sns.heatmap(confusion_matrix(test_labels, pred), annot=True)

```

```

Accuracy : 0.9441776710684273
              precision    recall  f1-score   support

      0         0.95        0.93        0.94        2393
      1         0.94        0.95        0.95        2605

 accuracy
macro avg         0.94        0.94        0.94        4998
weighted avg         0.94        0.94        0.94        4998

```

Out[81]: <AxesSubplot:>

